

Make sure to explain assumptions you make about the code you analyze. For instance, for pointer analysis, you may have made some assumptions about the aliasing information known about input parameters. Explain those assumptions and why they are reasonable.

Part of this project is to come up with a useful set of benchmarks on which to test and improve your analysis. Discuss why you chose those benchmarks, and what makes them interesting. If your implementation fails on some benchmarks (there's no shame in it!), then explain why and how the analysis might be improved.

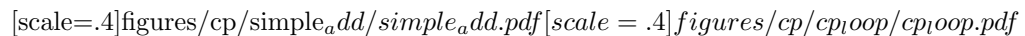
The difficulty of getting the various analyses to work properly on a piece of code is tightly coupled with the complexity of the underlying control flow graph. Pathologies that underly implementations of flow functions may not arise in straight-line programs but become painfully obvious when branches are introduced. Our benchmarks are designed to illustrate that our analyses are robust to non-trivial control flow structures.

Broadly speaking, we have three types of benchmark. The first type is straight-line programs, which introduce no branches in control structure. They are the easiest to handle, and our analyses are accordingly precise on them. Simple branching programs are the second type, and they introduce conditional branches into fold, but do not exhibit loops. They are slightly more challenging, but SSA makes them much easier to handle. Our final type of benchmark is looping programs. As their name suggests, they have loops, which makes precision quite difficult.

0.1 Benchmarks/Assumptions in Common

Because we have a common pool of benchmarks, we can list them here along with common assumptions on code. More specialized discussion of per-analysis benchmark goes below. Here, we can also introduce the Straight Line Program/Branching Program distinction.

0.2 Constant Propagation



0.3 Available Expressions

A straight line program where we can do CSE:

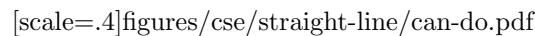


Figure 1: Available expression analysis of a straight-line program where CSE is possible

[scale=.4]figures/cse/straight-line/no-do.pdf [scale=.4]figures/cse/branch/can-do-cse.pdf [scale=.4]figures/cse/branch/no-do.pdf [scale=.4]figures/cse/loop/loop-can-do-cse.pdf [scale=.4]figures/cse/loop/loop-no-do-cse.pdf

0.4 Range Analysis

[scale=.4]figures/ra/simple_add/simple_add.pdf[scale=.4]figures/ra/simple_bbranch/simple_bbranch.pdf[scale=.4]figures/ra/loop/loop.pdf

0.5 Intra-Procedural Pointer Analysis

0.6 Range Checking

As per the assignment, we also implemented a range checking pass to make sure that array accesses remain inbound and raise warnings otherwise. This was fairly straightforward once the range analysis work was completed. To do this part, we simply ran our range analysis on the incoming function, then we scanned through the instructions for a `getelementptr` instruction. We then look at the lattice point at this instruction, query the index range, and raise a warning if it has a range that potentially falls outside of the range of the array.